

Handwritten Digit Classification

Brandon Park, Jungjae Park, Kyunghwan Min,
Minjae Lee, Minseo Jang

Tools

Language

Python

Random Forest

Tree Boosting

Single Layer Neural Network

Multilayer Neural Network

Data

MINIST Data set

Objective

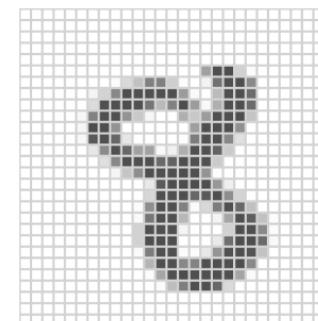
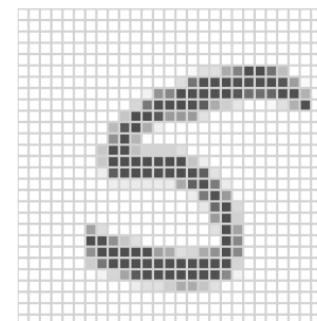
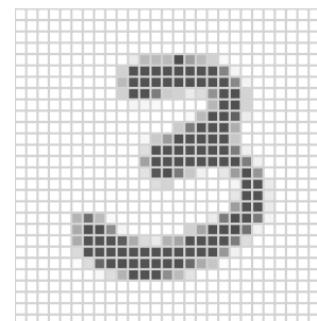
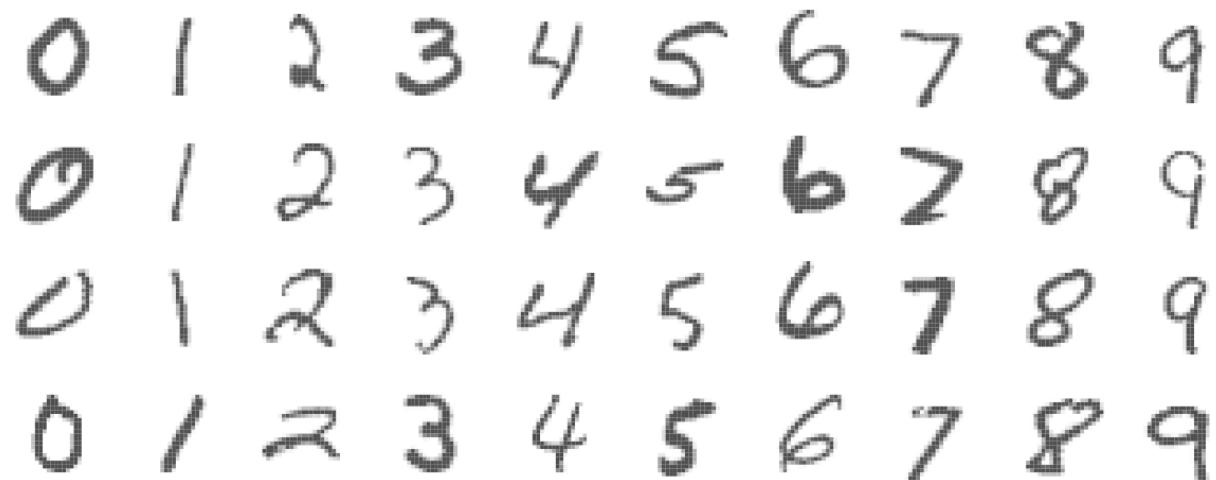
Compare the Efficiency

Time Cost

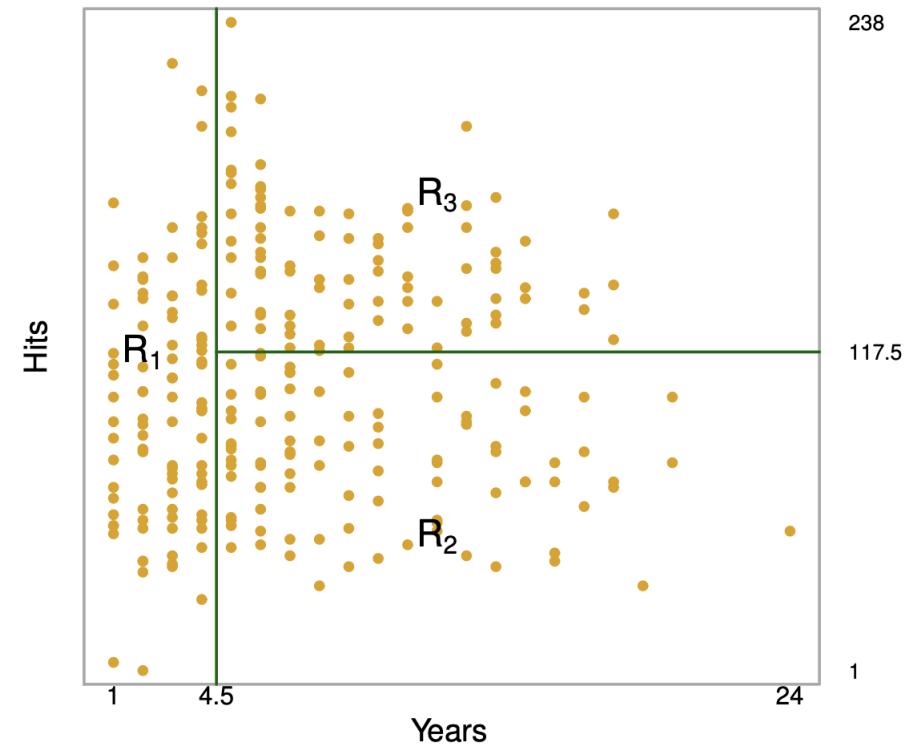
Accuracy

MNIST Data Set

- Training set:
60000 Images
- Test set: 10000 Images
- Each images with
 $28 \times 28 = 764$ pixels
- Each pixels have value
of 0 – 255



Trees



Building a Tree

- Splitting a Tree
 - Find boxes that minimize classification error rate

$$G = \sum_{k=1}^K \hat{p}_{mk} (1 - \hat{p}_{mk}),$$

$$D = - \sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}.$$

- Recursive binary splitting
 - Greedy, top-down approach

- Pruning a Tree
 - Overfitting -> poor test set performance
 - Option 1 : threshold
 - Option 2 : pruning – cost complexity pruning

$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T|$$

Random Forest

- 
- Problem 1 : **high variance**
 - Bootstrap : lowering variance by averaging a set of observations
 - Problem 2 : **correlation**
 - Same as bagging, at each split, a random sample of m predictors are only considered

Boosting

- Does not use bootstrapping
- Grows sequentially

1. Set $\hat{f}(x) = 0$ and $r_i = y_i$ for all i in the training set.
2. For $b = 1, 2, \dots, B$, repeat:
 - (a) Fit a tree \hat{f}^b with d splits ($d + 1$ terminal nodes) to the training data (X, r) .
 - (b) Update \hat{f} by adding in a shrunken version of the new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x).$$

- (c) Update the residuals,

$$r_i \leftarrow r_i - \lambda \hat{f}^b(x_i).$$

3. Output the boosted model,

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x).$$

Three parameters

B : Number of Trees

λ : Shrinkage parameter

d : Number of Splits

Result :

Random Forest

- Compared m values :
- 350 ($p/2$) vs 26 (\sqrt{p})
- 26 was better : less correlation
- Confusion matrix : strong diagonal

```
Number of predictors = 700
mtry1 (0.5p): 350
mtry2 (sqrt(p)): 26

--- Timing Random Forest (mtry1) ---
Time: 1396.43798494339 seconds

--- Timing Random Forest (mtry2) ---
Time: 432.87593960762024 seconds
OOB Error (0.5p): 0.03166666666666662
OOB Error (sqrt(p)): 0.02843333333333331
Best mtry selected = 26

--- Timing Final Random Forest ---
Time: 432.5059537887573 seconds
Test Misclassification Rate: 0.0277

Confusion Matrix:
[[ 970    0    0    0    0    0    3    2    1    3    1]
 [  0 1124    2    3    0    2    2    0    1    1    1]
 [  6    0 1000    5    3    0    4    8    6    0]
 [  0    0    9  976    0    5    0    9    9    2]
 [  1    0    2    0  958    0    4    0    2   15]
 [  3    0    0   10    3  862    6    1    5    2]
 [  6    3    0    0    3    3  939    0    4    0]
 [  1    2   17    1    1    0    0  994    1   11]
 [  3    0    5    7    2    5    3    3  935   11]
 [  5    5    2    9   10    2    1    4    6  965]]
```

Result : Tree Boosting

- Increasing depth led to lower error
 - Increasing number of trees led to lower error
-
- **Cannot see the effect of overfitting**

20000	3	0.0512	1926.30
20000	4	0.0372	2571.23
20000	5	0.0305	3405.74
20000	6	0.0282	4352.93
30000	1	0.1211	1341.37
30000	2	0.0651	2081.97
30000	3	0.0426	2870.53
30000	4	0.0314	3870.28
30000	5	0.0272	5155.69
30000	6	0.0261	6375.44

Best Parameters → Trees: 30000 | Depth: 6 | Validation Error: 0.0261

Overfitting

- Original task: **10-class digit classification (0–9)**
- Modified task: **Binary classification (digit 3 vs digit 8)**
- Reduces problem complexity and allows the model to reach high capacity more quickly

```
##### result #####
#Trees: 10000 | Depth: 4 | Validation Error Rate: 0.0129
#Trees: 10000 | Depth: 5 | Validation Error Rate: 0.0109
#Trees: 10000 | Depth: 6 | Validation Error Rate: 0.0121
#Trees: 20000 | Depth: 4 | Validation Error Rate: 0.01
#Trees: 20000 | Depth: 5 | Validation Error Rate: 0.0083
#Trees: 20000 | Depth: 6 | Validation Error Rate: 0.0104
#Trees: 30000 | Depth: 4 | Validation Error Rate: 0.0092
#Trees: 30000 | Depth: 5 | Validation Error Rate: 0.0083
#Trees: 30000 | Depth: 6 | Validation Error Rate: 0.0092
#Best Parameters → Trees: 20000 | Depth: 5 | Validation Error: 0.0083
#Test Misclassification Rate (Boosted Tree): 0.0055
#####
```

Single Layer Neural Network

- Hidden layer function : RELU
- Regularization : Dropout
- Searched over :
 - Neurons : 50, 100, 150, 200
 - Dropout : 0.2, 0.3, 0.4, 0.5
 - Batch size : 100, 200, 300
 - Learning rate : 0.01 (Adam)
 - Epoch : 5000
- **Best Parameters :**
Neurons: 150 | Dropout: 0.2 | Batch Size: 300
Validation Error: 3.19%
Test Error : 3.85%

Result : Single NN

- Increasing neurons generally improved validation error up to a point (capacity helps).
- Too much dropout (0.5) often worsened performance (too much regularization).
- Best NN achieved **3.85% test error**, but required **~31 minutes of grid search**, so tuning cost is significant.
- Compared to your best tree ensembles, this simple 1-hidden-layer network is competitive but not the top performer.

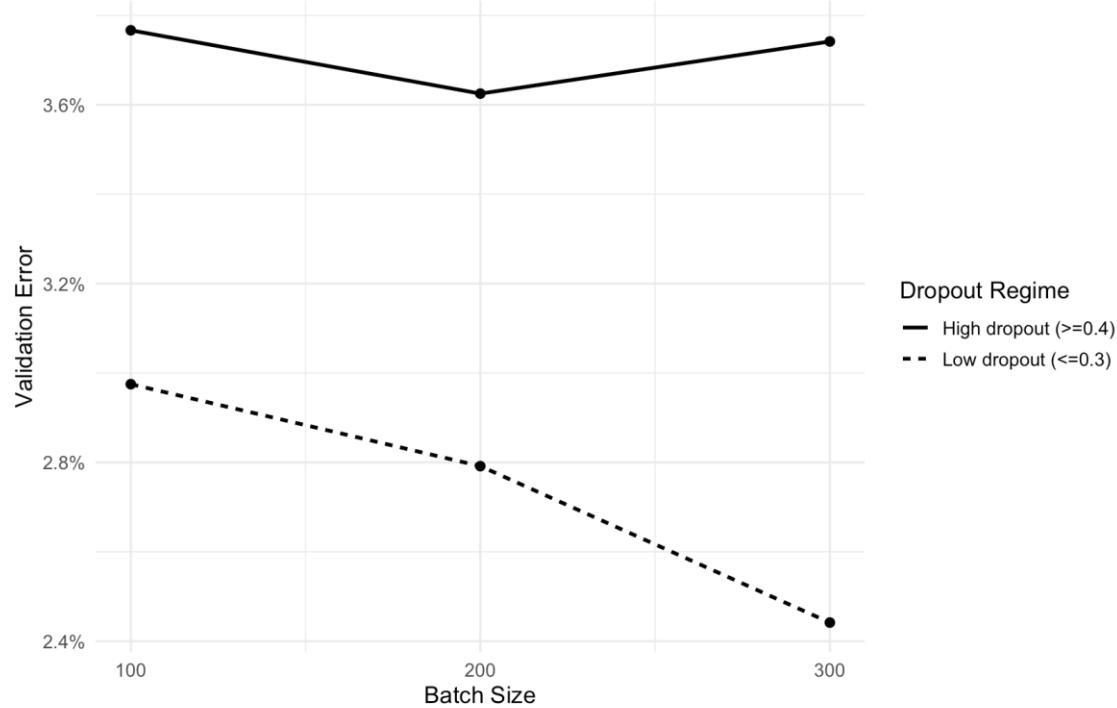
Multilayer Neural Network

- Learning rate: 0.001 (Adam)
- Epoch: 5000
- Searched for :
 - L1 (1st layer): 50, 100, 150, 200
 - L2 (2nd layer): 50, 100, 150, 200
 - Drop1: 0.2, 0.3, 0.4, 0.5
 - Drop2: 0.2, 0.3, 0.4, 0.5
 - Batch size: 100, 200, 300
- **Best Parameters:**
 - L1: 150 | L2: 50 | Drop1: 0.2 | Drop2: 0.2 | Batch Size: 300
- **Validation Error: 2.44%**
- **Test Error: 3.39%**

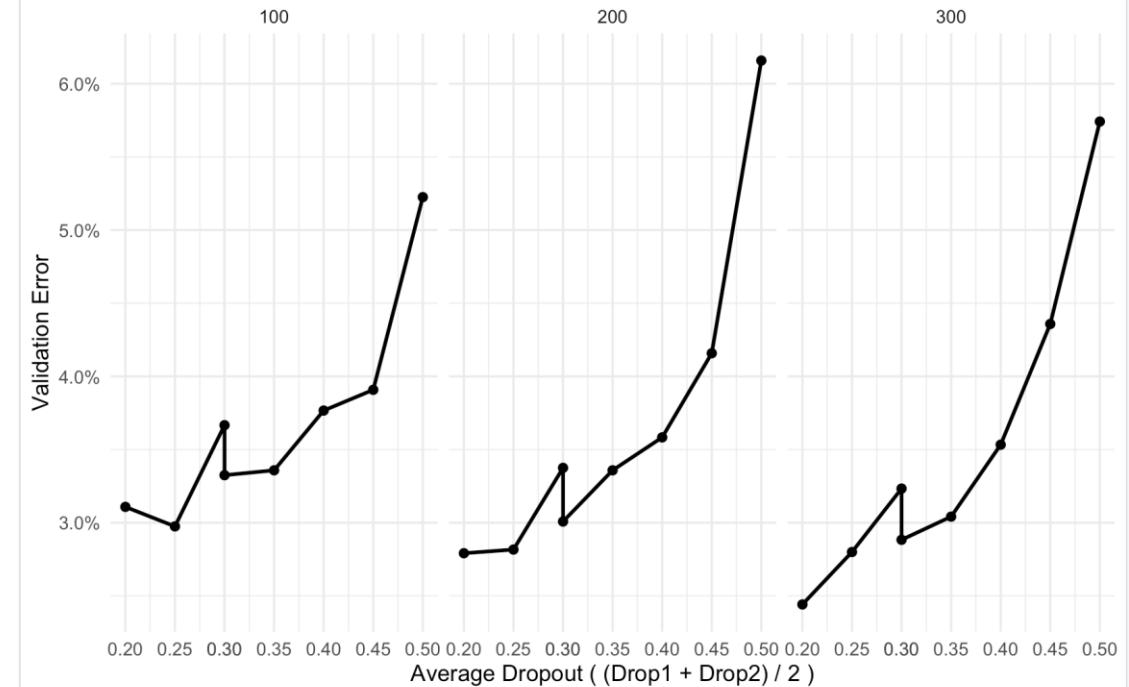
Multilayer Neural Network

- The second hidden layer mainly refines these features, so fewer units are sufficient.
- Low dropout provides mild regularization without removing too much information.
- Higher dropout values led to underfitting and worse validation performance.
- Larger batches produce more stable gradient updates when using Adam with a small learning rate.

Effect of Batch Size on Validation Error (Best-case per Batch)
Larger batches tend to give more stable/better results (especially under low dropout)



Effect of Dropout on Validation Error (Best-case per Dropout)
For each dropout level, we plot the minimum validation error across all architectures



Comparison with Single NN

- **Hidden unit allocation differs**
- Single-layer NN requires many units (150) in one layer because it performs feature extraction and classification together.
- Multi-layer NN distributes capacity across layers, reducing bias more efficiently.
- **Regularization strength is unchanged**
- Both models selected dropout = 0.2, showing that stronger regularization is unnecessary for MNIST.
- Over-regularization leads to underfitting in both cases.
- **Batch size remains the same**
- Batch size = 300 is optimal for training stability in both architectures.
- Optimization behavior does not change with depth.
- **Performance improvement source**
- The multi-layer NN improves accuracy by learning hierarchical representations.
- Gains are moderate, reflecting diminishing returns on this well-structured dataset.

References

- James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). *An introduction to statistical learning: With applications in R* (2nd ed.). Springer.
- LeCun, Y., Cortes, C., & Burges, C. J. C. (1998). **The MNIST database of handwritten digits.** <http://yann.lecun.com/exdb/mnist/>
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... Duchesnay, É. (2011). *Scikit-learn: Machine learning in Python*. *Journal of Machine Learning Research*, 12, 2825–2830.
- Chen, T., & Guestrin, C. (2016). *XGBoost: A scalable tree boosting system*. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* (pp. 785–794). ACM.
- Abadi, M., Agarwal, A., Barham, P., Brevdo, E., Chen, Z., Citro, C., ... Zheng, X. (2016). *TensorFlow: Large-scale machine learning on heterogeneous systems*. <https://www.tensorflow.org/>
- Chollet, F., et al. (2015). *Keras*. <https://keras.io>